

# Learning to Rank from Relevance Feedback

Peter Lubell-Doughtie

University of Amsterdam

`peter@heliod.com`

supervised by Maarten van Someren and Katja Hofmann

August 29th, 2011

# Learning to Rank from Relevance Feedback

- Search is inherently interactive, how can we use this to make search better?

# Learning to Rank from Relevance Feedback

- Search is inherently interactive, how can we use this to make search better?
- Our method uses judged documents for query expansion and then for learning to rank

# Learning to Rank from Relevance Feedback

- Search is inherently interactive, how can we use this to make search better?
- Our method uses judged documents for query expansion and then for learning to rank
- We improve search result ordering beyond that achieved when using only query expansion

# Learning to Rank from Relevance Feedback

- Search is inherently interactive, how can we use this to make search better?
- Our method uses judged documents for query expansion and then for learning to rank
- We improve search result ordering beyond that achieved when using only query expansion
- Success depends on how we extract features

# Learning to Rank from Relevance Feedback

- Search is inherently interactive, how can we use this to make search better?
- Our method uses judged documents for query expansion and then for learning to rank
- We improve search result ordering beyond that achieved when using only query expansion
- Success depends on how we extract features
- And on how we request document judgements

- Motivation
- Background
- Method
- Experiments and Results
- Conclusions

# Motivation: Learn from Users



# Motivation: Learn from Users

Search needs extra help when

# Motivation: Learn from Users

## Search needs extra help when

- Queries and documents contain ambiguous or domain specific terms

# Motivation: Learn from Users

## Search needs extra help when

- Queries and documents contain ambiguous or domain specific terms
- The user wants to retrieve many or all relevant documents

# Motivation: Learn from Users

## Search needs extra help when

- Queries and documents contain ambiguous or domain specific terms
- The user wants to retrieve many or all relevant documents
- There are multiple interactions with the search system

# Motivation: Learn from Users

## Search needs extra help when

- Queries and documents contain ambiguous or domain specific terms
- The user wants to retrieve many or all relevant documents
- There are multiple interactions with the search system

## Search is interactive, let's learn from this

# Motivation: Learn from Users

## Search needs extra help when

- Queries and documents contain ambiguous or domain specific terms
- The user wants to retrieve many or all relevant documents
- There are multiple interactions with the search system

## Search is interactive, let's learn from this

- Interaction is especially useful for difficult search tasks

# Motivation: Learn from Users

## Search needs extra help when

- Queries and documents contain ambiguous or domain specific terms
- The user wants to retrieve many or all relevant documents
- There are multiple interactions with the search system

## Search is interactive, let's learn from this

- Interaction is especially useful for difficult search tasks
- But providing explicit feedback is costly for the user

# Motivation: Learn from Users

## Search needs extra help when

- Queries and documents contain ambiguous or domain specific terms
- The user wants to retrieve many or all relevant documents
- There are multiple interactions with the search system

## Search is interactive, let's learn from this

- Interaction is especially useful for difficult search tasks
- But providing explicit feedback is costly for the user
- We must use the feedback we receive well



# Overarching Research Question

*How can a system that is aware of the interactive nature of search exploit this interactivity to better serve the user's needs?*

# Background: Learning to Rank from *Relevance Feedback*

## The Rocchio Algorithm and its Discontents

## The Rocchio Algorithm and its Discontents

- Combines an initial query, a positive contribution from relevant documents, and a negative contribution from non-relevant documents

## The Rocchio Algorithm and its Discontents

- Combines an initial query, a positive contribution from relevant documents, and a negative contribution from non-relevant documents
- Judged documents only affect ordering through the expanded query

## The Rocchio Algorithm and its Discontents

- Combines an initial query, a positive contribution from relevant documents, and a negative contribution from non-relevant documents
- Judged documents only affect ordering through the expanded query
- Negative feedback is often ignored because it is hard to use

## The Rocchio Algorithm and its Discontents

- Combines an initial query, a positive contribution from relevant documents, and a negative contribution from non-relevant documents
- Judged documents only affect ordering through the expanded query
- Negative feedback is often ignored because it is hard to use
- But negative feedback is found to be especially beneficial for difficult queries (Wang and Wang et al.)

# Background: Learning to Rank from *Relevance Feedback*

Our method addresses these challenges



# Background: Learning to Rank from *Relevance Feedback*

## Our method addresses these challenges

- To make better use of judged documents

# Background: Learning to Rank from *Relevance Feedback*

## Our method addresses these challenges

- To make better use of judged documents
- We use feedback in both learning to rank and query expansion

# Background: Learning to Rank from *Relevance Feedback*

## Our method addresses these challenges


- To make better use of judged documents
- We use feedback in both learning to rank and query expansion
- To avoid the problems of negative feedback

# Background: Learning to Rank from *Relevance Feedback*

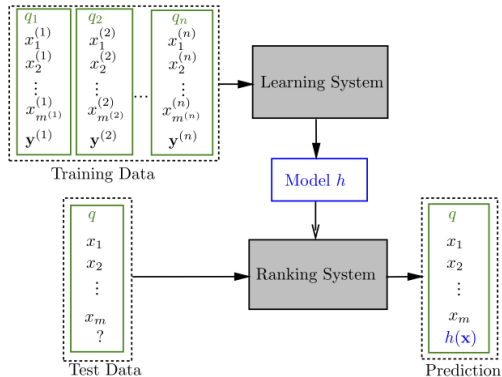
## Our method addresses these challenges

- To make better use of judged documents
- We use feedback in both learning to rank and query expansion
- To avoid the problems of negative feedback
- We use non-relevant documents in learning to rank, not query expansion

# Background: *Learning to Rank* from Relevance Feedback




# Background: *Learning to Rank* from Relevance Feedback




- Extract feature vectors from data

# Background: *Learning to Rank* from Relevance Feedback



- Extract feature vectors from data
- Use the training data to learn a model  $h$

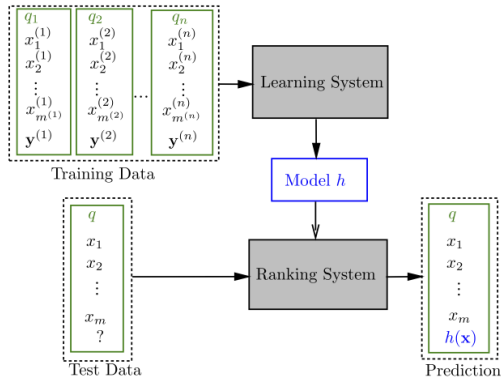
# Background: *Learning to Rank* from Relevance Feedback



- Extract feature vectors from data
- Use the training data to learn a model  $h$
- Use this model to predict the relevance of unlabeled test data



# Background: *Learning to Rank* from Relevance Feedback



- Extract feature vectors from data
- Use the training data to learn a model  $h$
- Use this model to predict the relevance of unlabeled test data
- Sort documents by their predicted relevance

# Background: *Learning to Rank* from Relevance Feedback

## Pairwise Learning to Rank

## Pairwise Learning to Rank

- Use document judgements to construct a preference pair set

## Pairwise Learning to Rank

- Use document judgements to construct a preference pair set
- $(i, j)$  in preference pairs iff document  $i$  preferred to document  $j$

# Background: *Learning to Rank* from Relevance Feedback

## Pairwise Learning to Rank

- Use document judgements to construct a preference pair set
- $(i, j)$  in preference pairs iff document  $i$  preferred to document  $j$
- E.g. if  $\mathbf{d}_1$  and  $\mathbf{d}_2$  are preferred to  $\mathbf{d}_3$ , then  $\mathcal{P} = \{(\mathbf{d}_1, \mathbf{d}_3), (\mathbf{d}_2, \mathbf{d}_3)\}$

# Background: *Learning to Rank* from Relevance Feedback

## Pairwise Learning to Rank

- Use document judgements to construct a preference pair set
- $(i, j)$  in preference pairs iff document  $i$  preferred to document  $j$
- E.g. if  $\mathbf{d}_1$  and  $\mathbf{d}_2$  are preferred to  $\mathbf{d}_3$ , then  $\mathcal{P} = \{(\mathbf{d}_1, \mathbf{d}_3), (\mathbf{d}_2, \mathbf{d}_3)\}$

## Minimize Incorrectly Ordered Pairs

# Background: *Learning to Rank* from Relevance Feedback

## Pairwise Learning to Rank

- Use document judgements to construct a preference pair set
- $(i, j)$  in preference pairs iff document  $i$  preferred to document  $j$
- E.g. if  $\mathbf{d}_1$  and  $\mathbf{d}_2$  are preferred to  $\mathbf{d}_3$ , then  $\mathcal{P} = \{(\mathbf{d}_1, \mathbf{d}_3), (\mathbf{d}_2, \mathbf{d}_3)\}$

## Minimize Incorrectly Ordered Pairs

- Using feature vectors for documents in the preference pair set

# Background: *Learning to Rank* from Relevance Feedback

## Pairwise Learning to Rank

- Use document judgements to construct a preference pair set
- $(i, j)$  in preference pairs iff document  $i$  preferred to document  $j$
- E.g. if  $\mathbf{d}_1$  and  $\mathbf{d}_2$  are preferred to  $\mathbf{d}_3$ , then  $\mathcal{P} = \{(\mathbf{d}_1, \mathbf{d}_3), (\mathbf{d}_2, \mathbf{d}_3)\}$

## Minimize Incorrectly Ordered Pairs

- Using feature vectors for documents in the preference pair set
- Calculate a weight vector over features



# Background: *Learning to Rank* from Relevance Feedback

## Pairwise Learning to Rank

- Use document judgements to construct a preference pair set
- $(i, j)$  in preference pairs iff document  $i$  preferred to document  $j$
- E.g. if  $\mathbf{d}_1$  and  $\mathbf{d}_2$  are preferred to  $\mathbf{d}_3$ , then  $\mathcal{P} = \{(\mathbf{d}_1, \mathbf{d}_3), (\mathbf{d}_2, \mathbf{d}_3)\}$

## Minimize Incorrectly Ordered Pairs

- Using feature vectors for documents in the preference pair set
- Calculate a weight vector over features
- Score unlabeled documents using this weight vector

# Background: A Difficult Recall Oriented Task

# Background: A Difficult Recall Oriented Task

e-Discovery and TREC Legal

# Background: A Difficult Recall Oriented Task

## e-Discovery and TREC Legal

- Ideally the search system returns *all* documents relevant to the query

# Background: A Difficult Recall Oriented Task

## e-Discovery and TREC Legal

- Ideally the search system returns *all* documents relevant to the query
- Ordered with the most relevant first

# Background: A Difficult Recall Oriented Task

## e-Discovery and TREC Legal

- Ideally the search system returns *all* documents relevant to the query
- Ordered with the most relevant first
- A hard task, the best systems have trouble with some queries

# Background: A Difficult Recall Oriented Task

## e-Discovery and TREC Legal

- Ideally the search system returns *all* documents relevant to the query
- Ordered with the most relevant first
- A hard task, the best systems have trouble with some queries
- Learning is important, previous TREC Legal participants' systems benefit substantially from incorporating feedback

# Background: A Difficult Recall Oriented Task

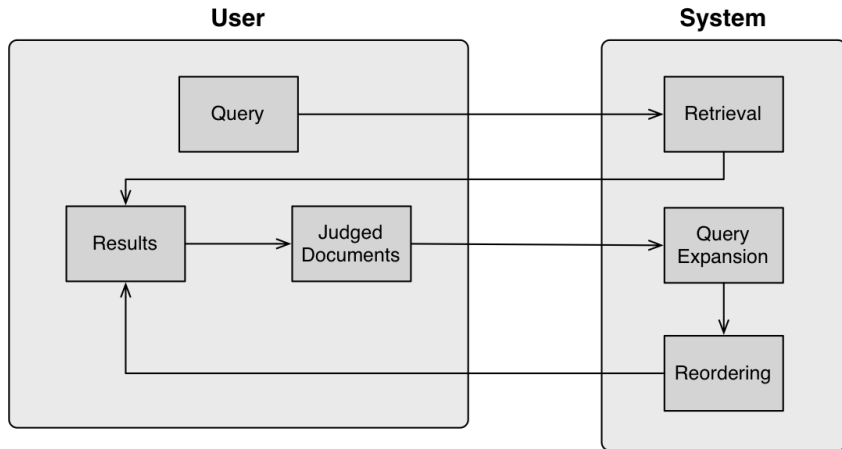
## e-Discovery and TREC Legal

- Ideally the search system returns *all* documents relevant to the query
  - Ordered with the most relevant first
- 
- A hard task, the best systems have trouble with some queries
  - Learning is important, previous TREC Legal participants' systems benefit substantially from incorporating feedback
  - We can simulate feedback by using already judged documents



# Method Overview

# Method Overview



# Open Questions

- What documents should the user judged?

# Open Questions

- What documents should the user judged?
- How do we represent documents for learning?

## Exploitative Sampling Strategy

# How to Request Judgements: Exploitative

## Exploitative Sampling Strategy

- Suppose we have a set of *seed* documents we expect are informative

## Exploitative Sampling Strategy

- Suppose we have a set of *seed* documents we expect are informative
- Sample these documents descending from the highest ranked



# How to Request Judgements: Exploitative

## Exploitative Sampling Strategy

- Suppose we have a set of *seed* documents we expect are informative
- Sample these documents descending from the highest ranked

## Example

- Suppose the returned documents are ordered:  $\mathbf{d}_1^* > \mathbf{d}_2 > \mathbf{d}_3^*$

# How to Request Judgements: Exploitative

## Exploitative Sampling Strategy

- Suppose we have a set of *seed* documents we expect are informative
- Sample these documents descending from the highest ranked

## Example

- Suppose the returned documents are ordered:  $\mathbf{d}_1^* > \mathbf{d}_2 > \mathbf{d}_3^*$
- Those marked \* are expected to be informative

# How to Request Judgements: Exploitative

## Exploitative Sampling Strategy

- Suppose we have a set of *seed* documents we expect are informative
- Sample these documents descending from the highest ranked

## Example

- Suppose the returned documents are ordered:  $\mathbf{d}_1^* > \mathbf{d}_2 > \mathbf{d}_3^*$
- Those marked \* are expected to be informative
- We would request judgements on  $\mathbf{d}_1$  and then  $\mathbf{d}_3$

# How to Request Judgements: Exploitative

## Exploitative Sampling Strategy

- Suppose we have a set of *seed* documents we expect are informative
- Sample these documents descending from the highest ranked

## Example

- Suppose the returned documents are ordered:  $\mathbf{d}_1^* > \mathbf{d}_2 > \mathbf{d}_3^*$
  - Those marked \* are expected to be informative
  - We would request judgements on  $\mathbf{d}_1$  and then  $\mathbf{d}_3$
- 
- Close to the way documents are judged in a real-life search system

# How to Request Judgements: Exploitative

## Exploitative Sampling Strategy

- Suppose we have a set of *seed* documents we expect are informative
- Sample these documents descending from the highest ranked

## Example

- Suppose the returned documents are ordered:  $\mathbf{d}_1^* > \mathbf{d}_2 > \mathbf{d}_3^*$
  - Those marked \* are expected to be informative
  - We would request judgements on  $\mathbf{d}_1$  and then  $\mathbf{d}_3$
- 
- Close to the way documents are judged in a real-life search system
  - Biases judged documents towards those ranked higher

# How to Request Judgements: Random Seeds

## Random Seed Sampling Strategy

# How to Request Judgements: Random Seeds

## Random Seed Sampling Strategy

- Randomly sample from the set of informative documents

# How to Request Judgements: Random Seeds

## Random Seed Sampling Strategy

- Randomly sample from the set of informative documents
- Same document set as in exploitative sampling



# How to Request Judgements: Random Seeds

## Random Seed Sampling Strategy

- Randomly sample from the set of informative documents
- Same document set as in exploitative sampling
- No rank based bias

# Open Questions

- What documents should the user judge?
  - 1 Exploitative
  - 2 Random

# Open Questions

- What documents should the user judge?
  - 1 Exploitative
  - 2 Random
- How do we represent documents for learning?

# Representing Documents: Cumulative Features

# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

## Example

# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

## Example

- Suppose there are three documents



# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

## Example

- Suppose there are three documents
- The retrieval scores before expansion are  $\langle 1, 2, 3 \rangle$

# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

## Example

- Suppose there are three documents
- The retrieval scores before expansion are  $\langle 1, 2, 3 \rangle$
- After judging one document on the first iteration

# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

## Example

- Suppose there are three documents
- The retrieval scores before expansion are  $\langle 1, 2, 3 \rangle$
- After judging one document on the first iteration
- The document profile is  $P = \{\mathbf{d}\}$  and the expansion term is  $t_1$

# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

## Example

- Suppose there are three documents
- The retrieval scores before expansion are  $\langle 1, 2, 3 \rangle$
- After judging one document on the first iteration
- The document profile is  $P = \{\mathbf{d}\}$  and the expansion term is  $t_1$
- Retrieval scores with the expanded query are  $\langle 1, 1, 3 \rangle$

# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

## Example

- Suppose there are three documents
- The retrieval scores before expansion are  $\langle 1, 2, 3 \rangle$
- After judging one document on the first iteration
- The document profile is  $P = \{\mathbf{d}\}$  and the expansion term is  $t_1$
- Retrieval scores with the expanded query are  $\langle 1, 1, 3 \rangle$
- We are only collecting features up to and including the first iteration

# Representing Documents: Cumulative Features

- Features are the retrieval scores for queries expanded with different sets of judged documents

## Example

- Suppose there are three documents
- The retrieval scores before expansion are  $\langle 1, 2, 3 \rangle$
- After judging one document on the first iteration
- The document profile is  $P = \{\mathbf{d}\}$  and the expansion term is  $t_1$
- Retrieval scores with the expanded query are  $\langle 1, 1, 3 \rangle$
- We are only collecting features up to and including the first iteration
- We are reordering the top two documents

## Example: Cumulative Features

The input to the *cumulative* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0, \emptyset, \langle 1, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 1, 1, 3 \rangle\} \\ \vdots \end{bmatrix},$$

## Example: Cumulative Features

The input to the *cumulative* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0, \emptyset, \langle 1, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 1, 1, 3 \rangle\} \\ \vdots \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 1, 2 \\ 1, 1 \end{bmatrix}.$$



## Example: Cumulative Features

The input to the *cumulative* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0, \emptyset, \langle 1, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 1, 1, 3 \rangle\} \\ \vdots \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 1, 2 \\ 1, 1 \end{bmatrix}.$$

- The top row are the retrieval scores without expansion

## Example: Cumulative Features

The input to the *cumulative* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0, \emptyset, \langle 1, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 1, 1, 3 \rangle\} \\ \vdots \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 1, 2 \\ 1, 1 \end{bmatrix}.$$

- The top row are the retrieval scores without expansion
- The bottom row are retrieval scores with expansion

## Example: Cumulative Features

The input to the *cumulative* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0, \emptyset, \langle 1, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 1, 1, 3 \rangle\} \\ \vdots \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 1, 2 \\ 1, 1 \end{bmatrix}.$$

- The top row are the retrieval scores without expansion
- The bottom row are retrieval scores with expansion
- The first column are features for the first document

## Example: Cumulative Features

The input to the *cumulative* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0, \emptyset, \langle 1, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 1, 1, 3 \rangle\} \\ \vdots \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 1, 2 \\ 1, 1 \end{bmatrix}.$$

- The top row are the retrieval scores without expansion
- The bottom row are retrieval scores with expansion
- The first column are features for the first document
- The second column are features for the second document

# Representing Documents: Constant Features

Features are the retrieval scores for the same set of judged documents but using different numbers of expansion terms

# Representing Documents: Constant Features

Features are the retrieval scores for the same set of judged documents but using different numbers of expansion terms

## Example

# Representing Documents: Constant Features

Features are the retrieval scores for the same set of judged documents but using different numbers of expansion terms

## Example

- Suppose there are three documents

# Representing Documents: Constant Features

Features are the retrieval scores for the same set of judged documents but using different numbers of expansion terms

## Example

- Suppose there are three documents
- The judged document profile is  $P = \{\mathbf{d}\}$



# Representing Documents: Constant Features

Features are the retrieval scores for the same set of judged documents but using different numbers of expansion terms

## Example

- Suppose there are three documents
- The judged document profile is  $P = \{\mathbf{d}\}$
- The first two expansion terms are  $\{t_1, t_2\}$

# Representing Documents: Constant Features

Features are the retrieval scores for the same set of judged documents but using different numbers of expansion terms

## Example

- Suppose there are three documents
- The judged document profile is  $P = \{\mathbf{d}\}$
- The first two expansion terms are  $\{t_1, t_2\}$
- The retrieval scores when expanding with  $t_1$  are  $\langle 2, 2, 3 \rangle$

# Representing Documents: Constant Features

Features are the retrieval scores for the same set of judged documents but using different numbers of expansion terms

## Example

- Suppose there are three documents
- The judged document profile is  $P = \{\mathbf{d}\}$
- The first two expansion terms are  $\{t_1, t_2\}$
- The retrieval scores when expanding with  $t_1$  are  $\langle 2, 2, 3 \rangle$
- The retrieval scores when expanding with  $t_1$  and  $t_2$  the retrieval scores are  $\langle 1, 2, 3 \rangle$

# Representing Documents: Constant Features

Features are the retrieval scores for the same set of judged documents but using different numbers of expansion terms

## Example

- Suppose there are three documents
- The judged document profile is  $P = \{\mathbf{d}\}$
- The first two expansion terms are  $\{t_1, t_2\}$
- The retrieval scores when expanding with  $t_1$  are  $\langle 2, 2, 3 \rangle$
- The retrieval scores when expanding with  $t_1$  and  $t_2$  the retrieval scores are  $\langle 1, 2, 3 \rangle$
- We are reordering the top two documents

## Example: Constant Features

Our input to the *constant* feature extractor  $\Phi$  will be:

$$\left[ \begin{array}{c} \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 2, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1 \circ t_2, \{\mathbf{d}\}, \langle 1, 2, 3 \rangle\} \end{array} \right],$$

## Example: Constant Features

Our input to the *constant* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 2, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1 \circ t_2, \{\mathbf{d}\}, \langle 1, 2, 3 \rangle\} \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 2, 2 \\ 1, 2 \end{bmatrix}.$$

## Example: Constant Features

Our input to the *constant* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 2, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1 \circ t_2, \{\mathbf{d}\}, \langle 1, 2, 3 \rangle\} \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 2, 2 \\ 1, 2 \end{bmatrix}.$$

- The top row are the retrieval scores for the query  $\mathbf{q}_0 \circ t_1$

## Example: Constant Features

Our input to the *constant* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 2, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1 \circ t_2, \{\mathbf{d}\}, \langle 1, 2, 3 \rangle\} \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 2, 2 \\ 1, 2 \end{bmatrix}.$$

- The top row are the retrieval scores for the query  $\mathbf{q}_0 \circ t_1$
- The bottom row are retrieval scores for the query  $\mathbf{q}_0 \circ t_1 \circ t_2$



## Example: Constant Features

Our input to the *constant* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 2, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1 \circ t_2, \{\mathbf{d}\}, \langle 1, 2, 3 \rangle\} \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 2, 2 \\ 1, 2 \end{bmatrix}.$$

- The top row are the retrieval scores for the query  $\mathbf{q}_0 \circ t_1$
- The bottom row are retrieval scores for the query  $\mathbf{q}_0 \circ t_1 \circ t_2$
- The first column are features for the first document

## Example: Constant Features

Our input to the *constant* feature extractor  $\Phi$  will be:

$$\begin{bmatrix} \{\mathbf{q}_0 \circ t_1, \{\mathbf{d}\}, \langle 2, 2, 3 \rangle\} \\ \{\mathbf{q}_0 \circ t_1 \circ t_2, \{\mathbf{d}\}, \langle 1, 2, 3 \rangle\} \end{bmatrix},$$

and the output features will be:

$$\begin{bmatrix} 2, 2 \\ 1, 2 \end{bmatrix}.$$

- The top row are the retrieval scores for the query  $\mathbf{q}_0 \circ t_1$
- The bottom row are retrieval scores for the query  $\mathbf{q}_0 \circ t_1 \circ t_2$
- The first column are features for the first document
- The second column are features for the second document

## Example: Term Frequency Features

- Features are the term frequency of the most common terms in judged documents and highly ranked documents

# Example: Term Frequency Features

- Features are the term frequency of the most common terms in judged documents and highly ranked documents

## Example

# Example: Term Frequency Features

- Features are the term frequency of the most common terms in judged documents and highly ranked documents

## Example

- Suppose there are three documents

## Example: Term Frequency Features

- Features are the term frequency of the most common terms in judged documents and highly ranked documents

### Example

- Suppose there are three documents
- Their retrieval scores are  $\langle 1, 2, 3 \rangle$

## Example: Term Frequency Features

- Features are the term frequency of the most common terms in judged documents and highly ranked documents

### Example

- Suppose there are three documents
- Their retrieval scores are  $\langle 1, 2, 3 \rangle$
- The term based features for the query are  $\mathbf{f}_q = [t_1, t_2]$

## Example: Term Frequency Features

- Features are the term frequency of the most common terms in judged documents and highly ranked documents

### Example

- Suppose there are three documents
- Their retrieval scores are  $\langle 1, 2, 3 \rangle$
- The term based features for the query are  $\mathbf{f}_q = [t_1, t_2]$
- For the first document the frequency of  $t_1$  is 1 and that of  $t_2$  is 2, i.e. its term frequency vector is  $[1, 2]$



## Example: Term Frequency Features

- Features are the term frequency of the most common terms in judged documents and highly ranked documents

### Example

- Suppose there are three documents
- Their retrieval scores are  $\langle 1, 2, 3 \rangle$
- The term based features for the query are  $\mathbf{f}_q = [t_1, t_2]$
- For the first document the frequency of  $t_1$  is 1 and that of  $t_2$  is 2, i.e. its term frequency vector is  $[1, 2]$
- For the second document the frequency of  $t_1$  is 1 and that of  $t_2$  is 1, i.e. its term frequency vector is  $[1, 1]$

## Example: Term Frequency Features

- Features are the term frequency of the most common terms in judged documents and highly ranked documents

### Example

- Suppose there are three documents
- Their retrieval scores are  $\langle 1, 2, 3 \rangle$
- The term based features for the query are  $\mathbf{f}_q = [t_1, t_2]$
- For the first document the frequency of  $t_1$  is 1 and that of  $t_2$  is 2, i.e. its term frequency vector is  $[1, 2]$
- For the second document the frequency of  $t_1$  is 1 and that of  $t_2$  is 1, i.e. its term frequency vector is  $[1, 1]$
- We are reordering the top two documents

## Example: Term Frequency Features

Our input to the term frequency feature extractor  $\Phi$  will be:

$$\left[ \{ \mathbf{q}, P_{\mathbf{q}}, \langle 1, 2, 3 \rangle \} \right],$$

## Example: Term Frequency Features

Our input to the term frequency feature extractor  $\Phi$  will be:

$$[\{\mathbf{q}, P_{\mathbf{q}}, \langle 1, 2, 3 \rangle\}] ,$$

and we define the output features as the term frequency vectors:

$$\begin{bmatrix} 1, 1 \\ 2, 1 \end{bmatrix} .$$

## Example: Term Frequency Features

Our input to the term frequency feature extractor  $\Phi$  will be:

$$[\{\mathbf{q}, P_{\mathbf{q}}, \langle 1, 2, 3 \rangle\}] ,$$

and we define the output features as the term frequency vectors:

$$\begin{bmatrix} 1, 1 \\ 2, 1 \end{bmatrix} .$$

- The top row are the term frequencies of  $t_1$

## Example: Term Frequency Features

Our input to the term frequency feature extractor  $\Phi$  will be:

$$[\{\mathbf{q}, P_{\mathbf{q}}, \langle 1, 2, 3 \rangle\}] ,$$

and we define the output features as the term frequency vectors:

$$\begin{bmatrix} 1, 1 \\ 2, 1 \end{bmatrix} .$$

- The top row are the term frequencies of  $t_1$
- The bottom row are the term frequencies of  $t_2$

## Example: Term Frequency Features

Our input to the term frequency feature extractor  $\Phi$  will be:

$$[\{\mathbf{q}, P_{\mathbf{q}}, \langle 1, 2, 3 \rangle\}] ,$$

and we define the output features as the term frequency vectors:

$$\begin{bmatrix} 1, 1 \\ 2, 1 \end{bmatrix} .$$

- The top row are the term frequencies of  $t_1$
- The bottom row are the term frequencies of  $t_2$
- The first column are features for the first document

## Example: Term Frequency Features

Our input to the term frequency feature extractor  $\Phi$  will be:

$$[\{\mathbf{q}, P_{\mathbf{q}}, \langle 1, 2, 3 \rangle\}] ,$$

and we define the output features as the term frequency vectors:

$$\begin{bmatrix} 1, 1 \\ 2, 1 \end{bmatrix} .$$

- The top row are the term frequencies of  $t_1$
- The bottom row are the term frequencies of  $t_2$
- The first column are features for the first document
- The second column are features for the second document



# Open Questions

# Open Questions

- What documents should the user judge?
  - ① Exploitative: choose highly ranked first
  - ② Random: choose at random

- What documents should the user judge?
  - ① Exploitative: choose highly ranked first
  - ② Random: choose at random
- How do we represent documents for learning?
  - ① Cumulative: retrieval scores for different document sets
  - ② Constant: retrieval scores with different numbers of expansion terms
  - ③ Term frequency: vector space model

# Learning to Rank improves Query Expansion

# Learning to Rank improves Query Expansion

## Random Seeds and Constant Features

Metric	Baseline	Expansion	Learning
MAP	0.0674	$\Delta$ 0.130	$\Delta$ <b>0.137</b>
NDCG	0.565	$\Delta$ 0.646	$\blacktriangle$ <b>0.652</b>

# Learning to Rank improves Query Expansion

## Random Seeds and Constant Features

Metric	Baseline	Expansion	Learning
MAP	0.0674	$\Delta$ 0.130	$\Delta$ <b>0.137</b>
NDCG	0.565	$\Delta$ 0.646	$\blacktriangle$ <b>0.652</b>

- *Baseline* is language model retrieval with the original query

# Learning to Rank improves Query Expansion

## Random Seeds and Constant Features

Metric	Baseline	Expansion	Learning
MAP	0.0674	$\Delta$ 0.130	$\Delta$ <b>0.137</b>
NDCG	0.565	$\Delta$ 0.646	$\blacktriangle$ <b>0.652</b>

- *Baseline* is language model retrieval with the original query
- *Expansion* is language model retrieval with the expanded query

# Learning to Rank improves Query Expansion

## Random Seeds and Constant Features

Metric	Baseline	Expansion	Learning
MAP	0.0674	$\Delta$ 0.130	$\Delta$ <b>0.137</b>
NDCG	0.565	$\Delta$ 0.646	$\blacktriangle$ <b>0.652</b>

- *Baseline* is language model retrieval with the original query
- *Expansion* is language model retrieval with the expanded query
- Learning from 20 relevant and 20 non-relevant documents per query



# Learning to Rank improves Query Expansion

## Random Seeds and Constant Features

Metric	Baseline	Expansion	Learning
MAP	0.0674	$\Delta$ 0.130	$\Delta$ <b>0.137</b>
NDCG	0.565	$\Delta$ 0.646	$\Delta$ <b>0.652</b>

- *Baseline* is language model retrieval with the original query
- *Expansion* is language model retrieval with the expanded query
- Learning from 20 relevant and 20 non-relevant documents per query
- Scores are averaged over all (8) queries

# Learning to Rank improves Query Expansion

## Random Seeds and Constant Features

Metric	Baseline	Expansion	Learning
MAP	0.0674	$\Delta$ 0.130	$\Delta$ <b>0.137</b>
NDCG	0.565	$\Delta$ 0.646	$\blacktriangle$ <b>0.652</b>

- *Baseline* is language model retrieval with the original query
- *Expansion* is language model retrieval with the expanded query
- Learning from 20 relevant and 20 non-relevant documents per query
- Scores are averaged over all (8) queries
- $\Delta$  are statistically significant at the 0.01 level and  $\blacktriangle$  at the 0.001 level

# Learning to Rank improves Query Expansion

*By combining traditional query expansion with learning to rank, a search system can use the interactive nature of search to better serve the user's needs.*

# How to Best Learn from Documents

# How to Best Learn from Documents

- The best method for representing documents is codependent with the best method for requesting document judgements

# How to Best Learn from Documents

- The best method for representing documents is codependent with the best method for requesting document judgements

## Experiment

- Compare evaluation scores when using different feature spaces


# How to Best Learn from Documents

- The best method for representing documents is codependent with the best method for requesting document judgements

## Experiment

- Compare evaluation scores when using different feature spaces
- Compare evaluation scores when using different strategies to choose documents for judgement

# Retrieval score features outperform term frequency features when using random sampling





# Why do retrieval score features outperform term frequency features when using random sampling?

# Why do retrieval score features outperform term frequency features when using random sampling?

- Random sampling judges documents at many different ranks

# Why do retrieval score features outperform term frequency features when using random sampling?

- Random sampling judges documents at many different ranks
- This provides an unbiased picture of the relationship between retrieval scores and document judgements

# Why do retrieval score features outperform term frequency features when using random sampling?

- Random sampling judges documents at many different ranks
- This provides an unbiased picture of the relationship between retrieval scores and document judgements
- Learning to rank with low ranked documents could increase the rank of other relevant but incorrectly low ranked documents

# Why do retrieval score features outperform term frequency features when using random sampling?

- Random sampling judges documents at many different ranks
- This provides an unbiased picture of the relationship between retrieval scores and document judgements
- Learning to rank with low ranked documents could increase the rank of other relevant but incorrectly low ranked documents
- This could significantly improve scores


# Why do retrieval score features outperform term frequency features when using random sampling?

- Random sampling judges documents at many different ranks
- This provides an unbiased picture of the relationship between retrieval scores and document judgements
- Learning to rank with low ranked documents could increase the rank of other relevant but incorrectly low ranked documents
- This could significantly improve scores
- Term frequency features only consider content and ignore rank

# Why do retrieval score features outperform term frequency features when using random sampling?

- Random sampling judges documents at many different ranks
- This provides an unbiased picture of the relationship between retrieval scores and document judgements
- Learning to rank with low ranked documents could increase the rank of other relevant but incorrectly low ranked documents
- This could significantly improve scores
- Term frequency features only consider content and ignore rank
- Content of low ranked documents may be uninformative

# Term frequency features outperform retrieval score features for exploitative sampling





# Why do term frequency features outperform retrieval score features when using exploitative sampling?

# Why do term frequency features outperform retrieval score features when using exploitative sampling?

- Exploitative sampling overrepresents highly ranked judged documents

# Why do term frequency features outperform retrieval score features when using exploitative sampling?

- Exploitative sampling overrepresents highly ranked judged documents
- With term frequency features the content of judged documents must be helpful for reordering

# Why do term frequency features outperform retrieval score features when using exploitative sampling?

- Exploitative sampling overrepresents highly ranked judged documents
- With term frequency features the content of judged documents must be helpful for reordering
- Our retrieval method orders documents based on the relevance of their terms

# Why do term frequency features outperform retrieval score features when using exploitative sampling?

- Exploitative sampling overrepresents highly ranked judged documents
- With term frequency features the content of judged documents must be helpful for reordering
- Our retrieval method orders documents based on the relevance of their terms
- Terms in highly ranked documents are likely to contain more informative terms than those from random documents

# Why do term frequency features outperform retrieval score features when using exploitative sampling?

- Exploitative sampling overrepresents highly ranked judged documents
- With term frequency features the content of judged documents must be helpful for reordering
- Our retrieval method orders documents based on the relevance of their terms
- Terms in highly ranked documents are likely to contain more informative terms than those from random documents
- $\text{frequency}(\text{top five terms in exploitative sampled documents}) = 5 \times \text{frequency}(\text{top five terms in random sampled documents})$

# Summary and Conclusions

- Our method combines learning to rank and query expansion

# Summary and Conclusions

- Our method combines learning to rank and query expansion
- This improves search result ordering beyond no expansion and beyond using only query expansion



# Summary and Conclusions

- Our method combines learning to rank and query expansion
- This improves search result ordering beyond no expansion and beyond using only query expansion
- Both constant and term frequency features can improve ordering

# Summary and Conclusions

- Our method combines learning to rank and query expansion
- This improves search result ordering beyond no expansion and beyond using only query expansion
- Both constant and term frequency features can improve ordering
- TREC Legal is not significantly different from other search task, such as ad-hoc web page search

# Summary and Conclusions

- Our method combines learning to rank and query expansion
- This improves search result ordering beyond no expansion and beyond using only query expansion
- Both constant and term frequency features can improve ordering
- TREC Legal is not significantly different from other search task, such as ad-hoc web page search
- Our method's improvement in scores may be a general result that will apply to other corpuses

- Design a better hybrid method

# Future Work

- Design a better hybrid method
- Apply to additional corpora

Questions?